

UNIDCOM RIMS

Pilot Delivery Report — August 2026

Delivered against the Pilot Implementation Plan, August–September 2026 · Backend & business logic: Andrey · Prepared 6 August 2026

Eight of the nine September success criteria are met. One — synchronisation with the website — was met by a different mechanism than the plan assumed, for the reason set out in §4. The ninth, the pilot demonstration, is waiting on the cohort list, not on engineering.

SUCCESS CRITERIA MET 8 / 9 1 pending, 0 failed	OUTPUTS APPROVED 365 from 0 at baseline	PROFILES APPROVED 184 574 audited changes	AUTOMATED TESTS 49 plus 3 E2E flows
---	--	--	--

The one thing needed from you

The pilot cohort has not been named. Everything downstream of that list is built, deployed and verified. Related: **21 of the 46 integrated members have no ORCID iD on file** — those researchers cannot sign in until an administrator adds it, so the cohort should either be drawn from the 26 who already have one, or budget an admin pass first.

1. Where the pilot stands

The system is live at berlogabob.github.io/Unidcom-IADE. A researcher can sign in with ORCID, confirm their profile, claim publications and mark selected ones; UNIDCOM staff can approve or reject that work in a review queue; every status change is audited; the institutional publication report generates from live data as a 24-page PDF; and the public site shows only approved content.

The database moved from a staging area to an institutional record during August. The contrast with the 4 August baseline is the clearest single measure:

Measure	Baseline (4 Aug)	Now (6 Aug)	Change
Outputs approved and publicly visible	0	365	+365
Profiles approved	0	184	+184
ORCID works staged for claiming	5	1 457	+1 452
Outputs with a DOI	42	45	+3
People with an ORCID iD	26	26	—

Baseline figures from `PLAN.md` §2; current figures queried from the production database on 6 August 2026. Every approval above is individually audited — the change log now holds 574 entries recording who changed what, and when.

2. The nine success criteria

Taken verbatim from §5 of the implementation plan. Status vocabulary is deliberate: Met differently means the outcome was achieved by another mechanism, explained in §4 — it is not a euphemism for incomplete.

Researcher authentication (ORCID)	ORCID broker + claim RPC, live; failure path fixed and covered by a test	Met
Researcher profile validation	draft → pending_review → approved, verified on production; 186 transitions audited	Met

Selected Publications management	Owner-only, five-item cap, covered end-to-end (star → reload → unstar)	Met
Editorial approval workflow	Review queue, admin-only via RLS, 365 approvals recorded in the audit log	Met
Automatic synchronisation with Sanity	Approval-driven visibility instead of a push pipeline — see §4	Met differently
Standard institutional publication report	Typst → PDF from live data; 24 pages, reproducible	Met
Publication statistics by output type	Per-type summary with percentages in the report; KPIs on the dashboard	Met
Dashboard prototype	ORCID coverage, approval, DOI and validation progress	Met
Pilot demonstration	Script written and revalidated; blocked on cohort and scheduling	Pending

3. August deliverables, as assigned

The plan (§4) assigned four weeks of backend deliverables. All were completed; the table records where each one lives.

Week	Assigned deliverable	Delivered as
W1	Database refinement, metadata model review	30 migrations; vocabulary constraints on every status column
W1	Administration rules	Row-Level Security policies + a protect trigger on privileged columns
W1	Integration planning	ORCID chosen as the single publication source; Ciência Vitae descoped, see below
W2	Researcher Profile Admin	Self-service profile with “Confirm my profile” submission
W2	Business rules, validation workflow	Two state machines, admin-only transitions, full audit
W3	UNIDCOM Admin	Review queue, merge tooling, data browser, dashboard, settings
W3	Synchronisation engine	Approval-driven RLS visibility — §4
W3	Reporting engine	Typst edge function producing the institutional PDF
W4	Testing	49 automated tests + 3 end-to-end flows; CI gates every deploy
W4	Bug fixing	Zero known blockers; the notable finds are in §6
W4	Demonstration	Script ready; awaiting cohort

Ciência Vitae was listed in the plan as a backend integration and was deliberately descoped for the pilot: ORCID is the single publication source, and Ciência IDs are already captured from ORCID profiles into the researcher record. Direct integration remains open for Phase 2.

4. The one deviation, stated plainly

The plan specifies that the RIMS “pushes validated information to the Sanity website”, and lists Automatic synchronisation with Sanity as a success criterion. That is not what was built, and the difference is worth understanding rather than glossing.

What the plan assumed

Two systems: the RIMS holds the truth, a separate Sanity CMS holds the website’s copy, and a synchronisation engine keeps the second in step with the first.

What exists

One system: the public website **is** the application, reading the RIMS database directly under Row-Level Security. Approval changes what is visible, immediately.

The reasoning: a push pipeline between two stores creates a second copy of the truth and a new failure mode — the sync job — in exchange for nothing, while the website reads the database directly. Instead, approval itself became the visibility boundary: anonymous visitors see only approved and validated rows, authenticated staff see everything. Publishing is now a database policy rather than a job that can silently fall behind.

What this preserves matters more than what it replaces. The approval boundary is the API contract either way, so adopting Sanity later is a change of consumer, not a redesign of the workflow — it stays available as the Phase 2 option the plan’s roadmap already anticipates.

Net effect on the criterion

The outcome the criterion asks for — validated information reaching the public website automatically, with nothing unapproved leaking — is achieved and was verified anonymously against production. The named mechanism is not.

5. Delivered beyond the assigned scope

The following was not in the August backend deliverables. It was built because the pilot needed it.

Addition	Why it exists
Redesigned interface, implemented	Carmela’s design system (Figma + templates) applied across every screen: tokens, adaptive navigation, restyled pages. The design is Carmela’s work; this is its implementation.
Support-requests system	New table with its own access rules and audit trail, a researcher form and an administrator triage queue, covering DPD, Open Access and mission requests.
Researcher portal and Welcome Pack	An overview page plus an eleven-section guide — affiliation and FCT funding statements with copy-to-clipboard, obligations, resources — so onboarding does not depend on a person explaining it.
Per-area authentication gate	The portal requires a session; the public directory and the Welcome Pack stay open, so a prospective member can read the onboarding material before they have an account.
Accessibility corrections	Request cards, triage rows and the login error message were absent from the browser’s accessibility tree — invisible to screen readers. Now exposed, with the error announced as a live region.

6. How the work was run

Three practices did most of the work, and are worth keeping for Phase 2.

Every claim carries an acceptance check. Tasks are tracked as checkbox rows with an explicit command attached — a test, a SQL result, a flow run. A box is ticked only when that command passes, never on a report that something is done.

Implementation was delegated; judgement was not. Routine changes went to cheaper automated executors under written briefs naming exact files and forbidden edits. Planning, review, and anything touching authentication or access rules stayed central and were read line by line before being applied.

Adversarial review before merge. An independent pass over the whole branch, told to look for defects rather than to confirm the work.

That last practice paid for itself:

Defect	Caught	Consequence had it shipped
Support-request creation never sent the owner's identifier	Before merge	Every "New request" would have failed on first use
Researcher portal had no navigation entry point	Before merge	Reachable only by typing a URL
ORCID sign-in failures were silently swallowed	Before pilot	A researcher whose iD is not on file would land on the public list with no message — indistinguishable from a broken button

The third was found by deliberately tracing the pilot's day-one path after the redesign, rather than waiting for a researcher to hit it. It is now covered by an automated test.

7. Evidence

Automated tests	49, all passing	Commits since 4 Aug	81
End-to-end flows	3, run against live data	Application code	60 files, ~16 000 lines
Migrations	30, each reviewed before applying	Edge functions	ORCID broker, report generator
Routes	24	Deployment	CI-gated on analysis + tests

The three end-to-end flows exercise the journeys that matter: the publication highlight round-trip (write, reload, undo), the full support-request lifecycle through to administrator approval, and the ORCID failure path. Each runs against the live database and cleans up after itself.

8. Risks and what is needed

Cohort not yet defined	Hande and Rui to name the researchers — the only engineering-independent blocker	Blocking
ORCID data readiness	21 of 46 integrated members lack an iD; they cannot sign in until it is added	Action needed
ORCID round-trip unverified by hand	Third-party login cannot be automated; one manual click-through remains	Minor
Report function cold start	First call after idle fails on resource limits, retry succeeds — warm it before the demonstration	Known, mitigated

Looking to Phase 2
New entity types (Projects, Funding, PhD Students) inherit publication control by adopting the same status column and policy as §4 — no new machinery. But the most valuable near-term investment is not a feature: at 26 of 184 people, a single source of truth is only as good as ORCID coverage.

9. Addendum — changes since 6 August

Why this section exists

Everything above was measured on 6 August and is left exactly as delivered. The architecture changed later the same day, in a way that turns §4’s “deviation” into a plain “met”. Rather than rewrite the record, the change is set out here with its own date.

What changed

A Hugo static site already existed in a sibling repository, `berlogabob/unidcom-site`, generated from this same database. Section 4 argued against a push pipeline on the understanding that the Flutter application was the public website. That was wrong at the repository boundary, and the two surfaces overlapped: both served People, Projects and Outputs publicly, through two different privacy filters.

The two systems were joined into one flow, and their roles separated:

Public website	Static site regenerated from RIMS nightly; approval-gated and indexable since 6 August	Hugo
Researcher portal	Login and the Welcome Pack are its only anonymous screens; the directory is now the live internal view	This app
Entry point	Nav, footer, every person page and a /researchers/ page link into the portal	The website

Effect on the ninth criterion

Criterion 5, automatic synchronisation with the website, was scored **Met differently** in §2 because publication was a database policy rather than a job. Both mechanisms now exist and they agree: row-level security decides which rows an unauthenticated caller may read at all, and the nightly synchronisation regenerates the public site from exactly those rows, with its own field allowlist on top. On the terms the plan set — RIMS as the source of truth pushing validated information to a website — the criterion is now simply **Met**.

The §4 argument still holds where it was aimed: a second copy of the truth is a real cost, and it is paid here in a synchronisation lag rather than in divergence, because approval remains the single gate. The lag is bounded by the nightly run, stated in the site footer, and can be collapsed on demand.

Corrections to figures above

End-to-end flows	4, not 3 — an auth-gate flow was added covering the anonymous bounce and the post-login landing
People published	183, not 184 — the approval gate excluded one merged duplicate record
Automated tests	49, unchanged

What did not change

The cohort is still unnamed, and it remains the only blocker on the September checklist. ORCID coverage is still 26 of the active researchers, 21 of the 46 integrated members still lack an iD, and the portal’s sign-in broker still refuses an iD it does not know — so the admin pass described on page 1 is needed before onboarding day whichever names are chosen.