

UNIDCOM RIMS

UX/UI Audit — researcher portal, 9 September 2026 (with same-day re-run)

Portal build 73259d4 (round 4, 8 Sep) · v1 pilot feature set · 86 screens, 7 flows · Prepared 9 September 2026

The pilot’s core loop works end to end — every journey that could be run passed. Two defects still block onboarding: the desktop navigation is invisible to keyboard and screen-reader users, and a person page can show, and let an admin edit, the wrong researcher. Neither is a design choice; both are a day’s fix with a test.

Release verdict: NOT READY at the first run — 2 blockers and 9 high-priority findings; READY WITH FIXES after the same-day fix round and re-run (§7). Sections 1–6 record the first run as found; §7 records what changed.

FINDINGS

43

2 sev-4 · 9 sev-3 · 26 sev-2 · 6 sev-1

TASK SUCCESS

100 %

7 of 7 runnable flows passed

BEST-PRACTICE SCORE

51 → 57

first run → re-run, of 38 items

The decision needed from you

Two deliberate August–September decisions carry a measurable usability bill: the one-to-one sidebar (25 rows, a third of them “Work in progress”) and the v1/v2 split that put DOI lookup behind the flag. Both are documented choices, not drift. This report prices them (§4, §5); revisiting them is a product call, not an engineering one.

1. What was audited, and how

The Flutter web portal at berlogabob.github.io/Unidcom-IADE, built from commit 73259d4 with the semantics tree on, in the three modes a user can be in — anonymous visitor, researcher, administrator — at desktop (1280 × 900) and phone (390 × 844) widths.

What	Source	Count
Screens captured (screenshot + accessibility tree each)	audit/2026-09-09-1136/screens.md	86
User journeys run end to end	flows-results.json — 2 not runnable: star an output (test account owns none), support requests (v2 only)	7 of 9
Provoked system states	cold start, wrong password, empty search, backend down at 3 s and 26 s	5
Review dimensions	Nielsen heuristics, UX laws, visual design, design system vs Figma, flows & states — plus orchestrator-verified findings	5 + 1
Best-practice checklist items scored	docs/research/2026-09-rims-best-practices.md, BP-01...38, each with a click-through or SQL test	38
Severity ≥ 3 claims re-tested before inclusion	2 corrected, 1 confirmed and promoted, 3 folded or dropped	6

Maestro, the repo’s E2E tool, could not drive Chromium on the audit machine (three attempts, driver never returned). The same journeys were run with Playwright; the scripts are committed under `audit/tools/` so the run is repeatable. The committed Maestro suite itself asserts text the build no longer shows (finding F-015).

2. The two blockers

Finding	Status	Evidence and cause
F-001 · Desktop navigation absent from the accessibility tree	Blocker	At ≥ 900 px none of the 25 sidebar rows, the mode switch or Sign out exists for assistive technology: 17 semantics nodes on the home page, 0 of them navigation. Verified after hover, keyboard focus, scroll, resize and reload. The phone drawer exposes every row correctly, so the defect is in the desktop layout branch.
F-002 · Person page keeps the previous researcher when only the id changes	Blocker	Two person URLs rendered byte-identical pages; after a hard reload the right person showed, and changing the id again kept the old one. Root cause: the page fetches once in <code>initState</code> and never on id change (<code>lib/public/person_page.dart:71</code>). An admin following a person-to-person link can edit the wrong record.

Both have a one-line fix and a one-test guard. Neither was visible to the existing tests, because no test renders the desktop shell at width or navigates person $\rightarrow$ person.

3. What a researcher will actually feel

Nine high-priority findings, in the order a researcher meets them.

ID	Where	What happens	Fix
F-007	Sidebar	25 rows in 6 groups; at 1280×900 only 11 fit before the fold, so Scientific Outputs is below it	Headers only by default; leaves on landing pages
F-008	Sidebar	The highlighted row is the wrong one on 6 of 25 pages (prefix match wins over exact)	Longest-match in <code>navSelected()</code>
F-009	My Outputs	A newly added output shows no pending tag — the researcher cannot tell what is approved	Reuse the existing status pill
F-005	Review queue	Reject fires with no confirmation, no reason, no feedback; Approve-all does confirm	Confirm + reason + Undo
F-003	Welcome pack	An anonymous visitor opening an M2 section hits a login wall instead of the promised redirect	Redirect to Getting started
F-004	Welcome pack	Signed-in and anonymous visitors get two different sidebars for the same page	One sidebar
F-006	404 page	Shows <code>GoException: no routes for location...</code> and drops the app shell	Friendly page inside the shell
F-010	Phone dashboard	KPI tiles clip at 390 px — key numbers unreadable	Wrap tiles
F-011	Logos page	Brand-colour swatch labels are illegible on their own swatch	Label outside the swatch

The remaining 32 findings (severity 2 and 1) are in `report.md`; the largest clusters are raw status enums shown to users (`draft`, `pending_review`, `a_confirmar`), mixed icon families including emoji, three accent colours on one KPI component, and a 20-second silent spinner before the (good) error message when the backend is unreachable.

4. Against best practice

The 38-item checklist distils what Pure, Converis, Elements, DSpace-CRIS and Haplo document, what CERIF, OpenAIRE and ORCID require, what FCT and CIÊNCIAVITAE ask of a unit, and the WCAG 2.2, GOV.UK and USWDS pattern rules that apply to forms and tables. Each item was tested by click-through or SQL on 9 September.

Met (9)		24 %
Partly met (21)		55 %
Not met (8)		21 %

Ahead of the field for its size. The data model — authorship as link rows with role and order, a unique DOI index, a full audit trail of 605 status changes — the approval gate that feeds a nightly, fail-closed public build (76 publications and 183 people on the site, matching the database), ORCID harvesting into a staging table for 26 researchers, and a real merge tool. Small units usually skip exactly these.

Behind every system surveyed. The entry experience:

Best practice	Status	What we have
BP-07 · Add output starts from a DOI/ORCID lookup	Not met	A blank six-field dialog; the Crossref client exists but sits behind the v2 flag, admin-only
BP-13 · Returned items carry a reviewer comment	Not met	Reject takes no reason; the item simply disappears from the researcher's list
BP-21 · Own list shows the workflow state as a text tag	Not met	No tag in My Outputs, although the pill component exists
BP-27 · Keyboard reach and 24 px targets	Not met	F-001; Reject is a 42 × 20 px text link
BP-16 · Duplicate check at entry (DOI, then ≈ 80 % title)	Partly	DOI is caught by the unique index; 8 duplicated normalised titles sit in the live table
BP-11/12 · Draft → submitted → validated, with withdraw	Partly	Outputs have no draft; a researcher's save is an immediate submission with no way back
BP-34/36 · FCT team export and data-quality counts	Partly	0 outputs flagged as representative; 1 466 unclaimed ORCID candidates not surfaced; CIÊNCIA ID coverage (25 of 184) not on the dashboard

5. Design criticism — built vs designed vs researched

Against the designer's templates. The tokens are faithful: navy sidebar, sand page, teal accent, 10 px radius, and the August contrast fix kept the brand teal while moving load-bearing text to the darker teal. The divergence is in components: template pills replaced by raw enum strings in five places, the template's icon set mixed with emoji in the Welcome pack, one KPI component in three accent colours, and a 404 page that drops the shell.

Against the research. Miller and Hick both argue for the six section headers as the visible menu, with leaves on the section landing pages — the model already has those pages. FCT's own guidance makes CIÊNCIA ID coverage and the five representative outputs the unit-level facts that matter; neither is on the dashboard. ORCID's display rules ask for the iD icon and the full URL; the profile shows a generic badge.

Drift, not decision. The sidebar semantics, the person-page refetch, the highlight logic, the stale Maestro suite and Reject-without-confirm were not chosen; each is a regression or omission that a test would have caught. That is the engineering message of this audit: the pilot needs a shell-width widget test and a navigate-A-to-B test more than it needs new features.

The public site is the stronger surface — fail-closed allowlists, ORCID links, sync green every night this week. Its open items are content (Portuguese bios under an English page, photos), not UX.

6. Recommended order of work

1. **F-001, F-002** — one day: fix both, pin each with a widget test.
2. **F-005, F-009, F-008** — one day: Reject confirm + reason + Undo; status pill in My Outputs; longest-match highlight.
3. **Sidebar** (F-007, F-023, F-035) — product decision with Rui, then half a day: headers-only default, mark WIP rows, fix the footer overlap.
4. **Welcome pack and 404** (F-003, F-004, F-006) — half a day.
5. **Phone** (F-010, F-036, F-037) — half a day.
6. **Regression suite** (F-015) — fix the six Maestro YAMLS and run them in CI, or adopt `audit/tools/flows.py`.
7. **BP-07 DOI-first entry, BP-16 title warning, BP-36 dashboard counts** — two days, the highest-value feature work before the demo.
8. **HEART instrumentation** — five events (`mode_chosen`, `profile_confirmed`, `output_added`, `candidate_claimed`, `review_decision`); three already exist in `change_log`.

7. Re-run after the fixes — same day

Round 5 fixed the two blockers and nine high-priority findings in ten pull requests (32–39 plus two follow-ups), each pinned by a widget test; the test suite went from 165 to 183. The audit tooling was then run again on the merged build (`audit/2026-09-09-1601`), same 86 screens, same seven flows, same five review passes, with the trend computed by (criterion, screen).

BLOCKERS
0
was 2

SEVERITY-WEIGHTED
SCORE
61
was 93 · target ≤ 60

FIXED / NEW /
REGRESSED
16 / 7 / 0
27 persisting, mostly by
decision

BEST-PRACTICE SCORE
57 / 100
was 51 · BP-21 met,
BP-13/27 partly

Metric	First run	Re-run
Findings by severity 4 / 3 / 2 / 1	2 / 9 / 26 / 6	0 / 1 / 25 / 8
Task success rate	100 %	100 %
Mean steps per flow	7.9	8.1
Defect density	0.50 per screen	0.41 per screen
Automated tests	165	183

Verified fixed on the live build: desktop sidebar exposed to assistive technology (42 semantics nodes on the researcher home, was 17); person page switches record on id change; anonymous M2 slug redirects; branded 404; correct sidebar highlight; status pill in My Outputs; Reject with confirmation, reason, Undo and the reason shown to the researcher; phone dashboard, review tabs and report table; logo labels; humanised status labels; chooser copy; sidebar footer; Maestro assertions.

Found by the re-run and fixed the same afternoon: the queue row printed “pending” twice after the pill change; the new 404 page rendered blank because it was wrapped in the app shell outside the router. Both now have tests — the argument for keeping this crawl in the routine.

Remaining severity 3 (one): the 25-row researcher sidebar (Miller/Hick), a product decision for Rui. The seven new findings are all severity 2 or 1: raw status chips on the researcher’s own profile header, three feedback patterns for one message type, the thin Personal Information page left by the E4 split, the v2

admin route landing on a queue instead of redirecting, cluster codes without expansion, and small alignment offsets.

Release verdict after the re-run

READY WITH FIXES

 — no blocker open; the one high-priority item is a design decision, not a defect.
The pilot cohort can be onboarded on this build.

8. Measurements for the next run

Metric	How measured	9 Sep 2026
Task success rate	7 runnable flows, all passed, 0 errors, 2 warnings	100 %
Mean steps per flow	tap/input count per journey	7.9
Findings by severity 4 / 3 / 2 / 1	findings.json	2 / 9 / 26 / 6
Severity - weighted score (Σ severity)	the number to drive down	93
Defect density	43 findings / 86 screens	0.50 per screen
Best-practice score	yes = 1, partial = ½, over 38 items	51 / 100
Researchers who can sign in	people.orcid filled — the adoption ceiling	26 of 184
Unclaimed ORCID candidates	output_candidates.status = 'pending'	1 466

Trend comparison is by (criterion, screen) pair between `findings.json` files; \$7 is the first such comparison. Rerun: `audit/tools/README.md`. Screenshots and accessibility trees are kept out of git — they contain researcher emails from the admin data browser — and live only in the audit machine's `audit/2026-09-09-1136/screens` and `hierarchy` folders.




Figure 1: Researcher home at 1280 × 900. The sidebar shown here has no accessibility nodes (F-001); “Add Scientific Output” is cut off behind the footer (F-035); the status tile shows the raw value `draft` (F-016).